Polymorphic Memory Resources - r2

Abstract

A significant impediment to effective memory management in C++ has been the inability to use allocators in non-generic contexts. In large software systems, most of the application program consists of non-generic procedural or object-oriented code that is compiled once and linked many times. Allocators in C++, however, have historically relied solely on compile-time polymorphism, and therefore have not been suitable for use in *vocabulary* types, which are passed through interfaces between separately-compiled modules, because the allocator type necessarily affects the type of the object that uses it. This proposal builds upon the improvements made to allocators in C++11 and describes a set of facilities for runtime polymorphic memory resources that interoperate with the existing compile-time polymorphic allocators. In addition, this proposal improves the interface and allocation semantics of some library classes, such as `std::function`, that use type erasure for allocators.

Contents

1 Proposal history

1.1 *Target*

The original version of this proposal ([N3525](#)) was first discussed in the Library Evolution Working Group during the April 2013 meeting of WG21 in Bristol, UK. A revised version of the proposal, N3726, was brought back to the LEWG at the September 2013 meeting in Chicago. A straw poll of the LEWG both in Bristol and in Chicago indicated strong support for the concepts in this proposal and a decision was made to target these ideas for inclusion a forthcoming Library Fundamentals Technical Specification (TS).

1.2 *Changes from N3816*

- Removed a number of tuning parameters from `pool_options`.

- Added a diagram of an example implementation of pool resources.

- Changed formal wording in the "Changes to classes in the standard" section to conform to recent decisions for how such changes should be described in the TS.

- Numerous small changes in response to a detailed technical review by the library working group.

1.3   Changes from N3726

- Added rationale in a few places to justify design choices. Removed guidance requests that have been already been addressed by the LEWG.

- Reorganized the formal wording into two sections: new classes and changes to standard classes. The intention is that both would go into the TS, even though the second section is not a pure extension.

- Moved all new classes and nested namespaces into the `std::experimental` namespace.

- Changed `allocator_resource` to use the public-non-virtual-function-calls-protected-virtual-function idiom.

- Clarified wording for various `do_allocate` and `do_deallocate` functions. Specifically, changed alignment requirements to reflect the conditionally-supported nature of superalignment. Also clarified that `do_deallocate` must be called on a block that was allocated from the same (or equal) resource.

- Made `resource_allocator_imp` constructor `explicit`.

- Added a `synchronized_pool_resource` class in addition to the previously-described `unsynchronized_pool_resource` class.

- Added some algorithmic description to the pool resources classes and added a few tuning parameters so that users can determine when their use is appropriate.

- Removed threshold for `monotonic_buffer_resource`.

- Added `noexcept` to `operator==` and `operator!=`. Changes `static const` to `constexpr`.

1.4   Changes from N3525

- Simplified alignment requirements for `memory_resource::allocate()`.

- Renamed the `polyalloc` namespace to `pmr` (Polymorphic Memory Resource).

- Simplified `new_delete_resource` and gave more leeway to the implementation.

- Added `null_memory_resource()` function.

- Borrowed some ideas from Mark Boyall's N3575 and mixed them with some ideas from Bloomberg's BSL project to yield the `monotonic_buffer_resource` and `unsynchronized_pool_resource` concrete manifestations of polymorphic memory resources.

- Specified allocator behavior for `promise` and `packaged_task`.

- There were some design changes proposed during discussion at the April 2013 meeting in Bristol. Although I elected not to make a number of those changes, I did investigate each of them and, for those ideas that were rejected, I added rationale for why they are the way they are.

- Wording improvements, especially in type-erased allocator section.

- Complete description of aliases for containers using polymorphic allocators.

2   Document Conventions

For the parts of this document that refer to changes in existing standard classes, section names and numbers are relative to the October 2013 Working Draft**, N3797.**

Existing working paper text is indented and shown in dark blue. Edits to the working paper are shown with ~~red strikeouts for deleted text~~ and <u>green underlining for inserted text</u> within the indented blue original text. When describing the addition of entirely new sections, the underlining is omitted for ease of reading.

Comments and rationale mixed in with the proposed wording appears as shaded text.

Requests for committee opinions and guidance appear with light (yellow) shading. It is expected that changes resulting from such guidance will be minor and will not delay acceptance of this proposal in the same meeting at which it is presented.

3   Motivation

Back in 2005, I argued in N1850 that the C++03 allocator model hindered the usability of allocators for managing memory use by containers and other objects that allocate memory. Although N1850 conflated them, the proposals in that paper could be broken down into two separate principles:

1. The allocator used to construct a container should also be used to construct the elements within that container.

2. An object's type should be independent of the allocator it uses to obtain memory.

In subsequent proposals, these principles were separated. The first principle eventually became known as the scoped allocator model and is embodied in the `scoped_allocator_adaptor` template in Section [allocator.adaptor] (20.12) of the 2011 standard (and the same section of the current WP).

Unfortunately, creating a scoped allocator model that was compatible with C++03 and acceptable to the committee, as well as fixing other flaws in the allocator section of the standard, proved a time-consuming task, and library changes implementing the second principle were not proposed in time for standardization in 2011.

This paper proposes new library facilities to address the second principle. Section 4.3 of N1850 (excerpted in the appendix of this paper) gives a detailed description of why it is undesirable to specify allocators as class template parameters. Key among the problems of allocator template parameters is that they inhibit the use of

vocabulary types by altering the type of specializations that would otherwise be the same. For example, `std::basic_string<char, char_traits<char>, Alloc1<char>>` and `std::basic_string<char, char_traits<char>, Alloc2<char>>` are different types in C++ even though they are both string types capable of representating the same set of (mathematical) values.

Some new vocabulary types introduced into the 2011 standard, including `function`, `promise`, and `future` use *type erasure* (see [jsmith]) as a way to get the benefits of allocators without the allocator contaminating their type. Type erasure is a powerful technique, but has its own flaws, such as that the allocators can be propagated outside of the scope in which they are valid and also that there is no way to query an object for its type-erased allocator. More importantly, even if type erasure were a completely general solution, it cannot be applied to existing container classes because they would break backwards compatibility with the existing interfaces and binary compatibility with existing implementations. Moreover, even for programmers creating their own classes, unconstrained by existing usage, type-erasure is a relatively complex and time-consuming technique and requires the creation of a polymorphic class hierarchy much like the `memory_resource` and `resource_adaptor` class hierarchy proposed for standardization below. Given that type erasure is expensive to implement and not general even when it is feasible, we must look to other solutions.

Fortunately, the changes to the allocator model made in 2011 (especially full support for stateful allocators and scoped allocators) make this problem with allocators relatively easy to solve in a more general way. The solution presented in this paper is to create a uniform memory allocation base class, `memory_resource`, suitable for use by template and non-template classes alike, and single allocator template, `polymorphic_allocator` that wraps a pointer to a `memory_resource` and which can be used ubiquitously for instantiating containers. The `polymorphic_allocator` will, as its name suggests, have polymorphic runtime behavior. Thus objects of the same type can have different effective allocators, achieving the goal of making an object's type independent of the allocator it uses to obtain memory, and thereby allowing them to be interoperable when used with precompiled libraries.

4   Usage Example

Suppose we are processing a series of shopping lists, where a shopping list is a container of strings, and storing them in a collection (a list) of shopping lists. Each shopping list being processed uses a bounded amount of memory that is needed for a short period of time, while the collection of shopping lists uses an unbounded amount of memory and will exist for a longer period of time. For efficiency, we can use a more time-efficient memory allocator based on a finite buffer for the temporary shopping lists. However, this time-efficient allocator is not appropriate for the longer lived collection of shopping lists. This example shows how those temporary shopping lists, using a time-efficient allocator, can be used to populate the long lived collection of shopping lists, using a general purpose allocator, something that would be annoyingly difficult without the polymorphic allocators in this proposal.

First, we define a class, `ShoppingList`, that contains a vector of strings. It is not a template, so it has no `Allocator` template argument. Instead, it uses `memory_resource` as a way to allow clients to control its memory allocation:

```cpp
#include <polymorphic_allocator>
#include <vector>
#include <string>

class ShoppingList {
    // Define a vector of strings using polymorphic allocators.  polymorphic_allocator is scoped,
    // so every element of the vector will use the same allocator as the vector itself.
    typedef std::pmr::string   string_type;
    typedef std::pmr::vector<string_type> strvec_type;

    strvec_type m_strvec;

  public:
    // This type makes uses_allocator<ShoppingList, memory_resource*>::value true.
    typedef std::pmr::memory_resource* allocator_type;

    // Construct with optional memory_resource.  If alloc is not specified, uses pmr::get_default_resource().
    ShoppingList(allocator_type alloc = nullptr)
      : m_strvec(alloc) { }

    // Copy construct with optional memory_resource.
    //  If alloc is not specified, uses pmr::get_default_resource().
    ShoppingList(const ShoppingList& other) = default;
    ShoppingList(std::allocator_arg_t, allocator_type a,
                 const ShoppingList& other)
      : m_strvec(other, a) { }

    allocator_type get_allocator() const
        { return m_strvec.get_allocator().resource(); }

    void add_item(const string_type& item){ m_strvec.push_back(item); }
    ...
};

bool operator==(const ShoppingList &a, const ShoppingList &b);
```

There was some discussion in LEWG as to whether it was appropriate to use `allocator_type` as an alias for something that is not, strictly speaking, an allocator. At the time, I sympathized with this objection and set out to see what the ripple effect would be if a different `typedef` name were chosen in cases where a class uses `memory_resource` directly. Unfortunately, the ripple effect is too large, in my opinion, to justify this change. In particular, every class function or constructor that propagates its allocator to a member or element would need to be reworded to use argument names like `allocator_or_resource` and descriptions with duplicate wording based on whether an allocator or resource pointer were passed in. In effect, we would be undoing in English what we so carefully created in the interface, which is the nearly complete interchangeability of allocators and memory resource pointers.

Next, we create an allocator resource, `FixedBufferResource`, that allocates memory from a fixed-size buffer supplied at construction. The `FixedBufferResource` is not responsible for reclaiming this externally managed buffer, and consequently its `deallocate` method and destructor are no-ops. This makes allocations and deallocations very fast, and is useful when building up an object of a bounded size that will be destroyed all at once (such as one of the short lived shopping lists in this example).

```cpp
class FixedBufferResource : public std::pmr::memory_resource
{
    void*       m_next_alloc;
    std::size_t m_remaining;

  public:
    FixedBufferResource(void* buffer, std::size_t size)
      : m_next_alloc(buffer), m_remaining(size) { }

  protected:
    virtual void* do_allocate(std::size_t sz, std::size_t alignment)
    {
        if (std::align(alignment, sz, m_next_alloc, m_remaining))
        {
            void* ret = m_next_alloc;
            m_next_alloc = static_cast<char*>(m_next_alloc) + sz;
            return ret;
        }
        else
            throw std::bad_alloc();
    }

    virtual void do_deallocate(void*, std::size_t, std::size_t) { }

    virtual bool do_is_equal(std::pmr::memory_resource& other) const
        noexcept
    {
        return this == &other;
    }
};
```

Now, we use the `ShoppingList` and `FixedBufferResource` defined above to demonstrate processing a short-lived shopping list into a collection of shopping lists. We define a collection of shopping lists, `folder`, that will use the default allocator. The temporary shopping list `temporaryShoppingList` will use the `FixedBufferResource` to allocator memory, since the items being added to the list are of a fixed size.

Note that the memory-resource library is designed so that the `ShoppingList` constructor accepts a *pointer* to a `memory_resource` rather than a *reference* to a `memory_resource`. It was noted that one common practice is to use references rather than pointers in situations where a null pointer is out of contract. However, there is a more compelling practice of avoiding constructors that take objects by reference and store their addresses. We also want to avoid passing non-const

references, as that, too, is usually considered bad practice (except in overloaded operators).

```cpp
std::pmr::list<ShoppingList> folder;    // Default allocator resource
{
    char buffer[1024];
    FixedBufferResource buf_rsrc(&buffer, 1024);
    ShoppingList temporaryShoppingList(&buf_rsrc);
    assert(&buf_rsrc == temporaryShoppingList.get_allocator());

    temporaryShoppingList.add_item("salt");
    temporaryShoppingList.add_item("pepper");

    if (processShoppingList(temporaryShoppingList)) {
        folder.push_back(temporaryShoppingList);
        assert(std::pmr::get_default_resource() ==
               folder.back().get_allocator());
    }

    // temporaryShoppingList, buf_rsrc, and buffer go out of scope
}
```

Notice that the shopping lists within `folder` use the default allocator resource whereas the shopping list `temporaryShoppingList` uses the short-lived but very fast `buf_rsrc`. Despite using different allocators, you can insert `temporaryShoppingList` into `folder` because they have the same `ShoppingList` type. Also, while `ShoppingList` uses `memory_resource` directly, `std::pmr::list`, `std::pmr::vector`, and `std::pmr::string` all use `polymorphic_allocator`. The resource passed to the `ShoppingList` constructor is propagated to the vector and each string within that `ShoppingList`. Similarly, the resource used to construct `folder` is propagated to the constructors of the `ShoppingList`s that are inserted into the list (and to the strings within those `ShoppingList`s). The `polymorphic_allocator` template is designed to be almost interchangeable with a pointer to `memory_resource`, thus producing a "bridge" between the template-policy style of allocator and the polymorphic-base-class style of allocator.

5   Summary of Proposal

5.1   Namespace `std::pmr`

All new components introduced in this proposal are in a new namespace, `pmr`, nested within namespace `std`.

The name, `pmr`, and all other identifiers introduced in this proposal are subject to change. If this proposal is accepted, we can have the bicycle-shed discussion of names. If you think of a better name, send a suggestion to the email address at the top of this paper.

5.2 Abstract base class `memory_resource`

An abstract base class, `memory_resource`, describes a memory resource from which blocks can be allocated and deallocated. It provides functions `allocate()`, `deallocate()`, and `is_equal()`, which call pure virtual functions `do_allocate()`, `do_deallocate()`, and `do_is_equal()`, respectively. Derived classes of `memory_resource` contain the machinery for actually allocating and deallocating memory. Note that `memory_resource`, not being a template, operates at the level of raw bytes rather than objects. The caller is responsible for constructing objects into the allocated memory and destroying the objects before deallocating the memory.

5.3 Class Template `polymorphic_allocator<T>`

An instance of `polymorphic_allocator<T>` is a wrapper around a `memory_resource` pointer that gives it a C++11 allocator interface. It is this adaptor that achieves the goal of separating an object's type from its allocator, especially for existing templates that have an allocator template parameter. Two objects x and y of type `list<int, polymorphic_allocator<int>>` have the same type, but may use different memory resources.

Polymorphic allocators use scoped allocator semantics. Thus, a container containing other containers or strings can be built to use the same memory resource throughout if polymorphic allocators are used ubiquitously.

5.4 Aliases for container classes

There would be an alias in the `pmr` namespace for each standard container (except `array`). The alias would not take an allocator parameter but instead would use `polymorphic_allocator<T>` as the allocator. For example, the `<vector>` header would contain the following declaration:

```
namespace std {
namespace pmr {

template <class T>
  using vector<T> = std::vector<T, polymorphic_allocator<T>>;

} // namespace pmr
} // namespace std
```

Thus, `std::pmr::vector<int>` would be a vector that uses a polymorphic allocator. Consistent use of his aliases would allow `std::pmr::vector<int>` to be used as a vocabulary type, interoperable with all other instances of `std::pmr::vector<int>`.

Within the LEWG, there was extensive discussion of the desirability of creating same-name aliases within a nested namespace. Proponents argued that the name `std::pmr::vector` would be cleaner and better accepted than `pmr_vector` or `std::pmr::pmr_vector`. Opponents claimed that users were likely to run into ambiguities if both `using std;` and `using std::pmr;` were present (though such an

ambiguity would be noisy and thus easy to fix). A straw poll was strongly in favor of leaving the aliases as proposed here (and warning users not to put `using std::pmr` in their code).

5.5   Class template `resource_adaptor<Alloc>`

An instance of `resource_adaptor<Alloc>` is a wrapper around a C++11 allocator type that gives it a `memory_resource` interface. In a sense, it is the complementary adaptor to `polymorphic_allocator<T>`. The adapted allocator, `Alloc`, is required to use normal (raw) pointers, rather than shared-memory pointers or pointers to some other kind of weird memory. (I have floated the term, *Euclidean Allocator*, to describe allocators such as these ☺.) The `resource_adaptor` template is actually an alias template designed such that `resource_adaptor<X<T>>` and `resource_adaptor<X<U>>` are the same type for all parameters `T` and `U`.

5.6   Function `new_delete_resource()`

Returns a pointer to a memory resource that forwards all calls to `allocate()` and `deallocate()` to global `operator new()` and `operator delete()`, respectively. Every call to this function returns the same value. Since the resource is stateless, all instances of such memory resources would be equivalent and there is never a need for more than one instance in a program.

5.7   Function `null_memory_resource()`

Returns a pointer to a memory resource that always fails with a `bad_alloc` exception when `allocate()` is called. This function is useful for setting the end of a *chain* of memory resource, where one memory resource depends on another. In cases where the first memory resource is not expected to exhaust its own pool of memory, the null memory resource can be used to avoid accidentally allocating memory from the heap. This function is also useful for testing, in situations such as the small-object optimization, where an allocator must be supplied, but is not expected to be used.

5.8   Functions `get_default_resource()` and `set_default_resource()`

Namespace-scoped functions `get_default_resource()` and `set_default_resource()` are used to get and set a specific memory resource to be used by certain classes when an explicit resource is not specified to the class's constructor. The ability to change the default resource used when constructing an object is extremely useful for testing and can also be useful for other purposes such as preventing DoS attacks by limiting the maximum size of an allocation.

If `set_default_resource()` is never called, the "default default" memory resource is `new_delete_resource()`.

5.9   Standard memory resources

A new library facility for using different types of allocators is useful only to the extent that such allocators actually exist. This proposal, therefore, includes a few memory

resource classes that have broad usefulness in our experience. In the future, we may propose additional resource classes for standardization, including a resource for testing the memory allocation behavior of allocator-aware classes.

5.9.1 Classes `synchronized_pool_resource` and `unsynchronized_pool_resource`

The `synchronized_pool_resource` and `unsynchronized_pool_resource` classes are general-purpose resources that *own* the allocated storage and free it on destruction, even if `deallocate` is not called for some or all of the allocated blocks. Efficiency is obtained by allocating memory in chunks from an "upstream" allocator (often the default allocator) and by maximizing storage locality among separate allocations. A logical data structure would be a set of object pools, but the actual choice of data structure and algorithm is left to the QOI.

5.9.2 Class `monotonic_buffer_resource`

The `monotonic_buffer_resource` class is designed for very fast memory allocations in situations where memory is used to build up a few objects and then is released all at once when those objects go out of scope. Like `unsynchronized_pool_resource`, it owns its memory and it is intended for single-threaded operation. The "monotonic" in its name refers to the fact that its use of memory increases monotonically because its `deallocate()` member is a no-op. By ignoring deallocation calls, this type of memory resource can use extremely simple data structures that do not require keeping track of individual allocated blocks. In addition, the user can provide it an initial buffer from which to allocate memory. In many applications, this buffer can reside on the stack, providing even more efficient allocation for small amounts of memory.

A particularly good use for a `monotonic_buffer_resource` is to provide memory for a local variable of container or string type. For example, the following code concatenates two strings, looks for the word "hello" in the concatenated string, and then discards the concatenated string after the word is found or not found. The concatenated string is expected to be no more than 80 bytes long, so the code is optimized for these short strings using a small `monotonic_buffer_resource` (but will still work, using the default allocator as a backup resource, if the concatenated string is over 80 bytes long):

```
bool find_hello(const std::pmr::string s1, const std::pmr::string s2)
{
    char buffer[80];
    monotonic_buffer_resource m(buffer, 80);
    std::pmr::string s(&m);
    s.reserve(s1.length() + s2.length());
    s += s1;
    s += s2;
    return s.find("hello") != pmr::string::npos;
    // s goes out of scope, then m and buffer go out of scope
}
```

5.10 Idiom for type-Erased Allocators

Type-erased allocators, which are used by `std::function`, `std::promise`, and `std::packaged_task` are already implemented internally using polymorphic wrappers. In this proposal, the implicit use of polymorphic wrappers is made explicit (reified). When one of these types is constructed, the caller may supply either a C++11 allocator or a pointer to `memory_resource`. A new member function, `get_memory_resource()` will return a pointer to the memory resource or, in the case where a C++11 allocator was provided at construction, a pointer to a `resource_adaptor` containing the original allocator. This pointer can be used to create other objects using the same allocator. If no allocator or resource was provided at construction, the value of `get_default_resource()` is used. To complete the idiom, classes that use type-erased allocators will declare

```
typedef erased_type allocator_type;
```

indicating that the class uses allocators, but that the allocator is type-erased. (`erased_type` is an empty class that exists solely for this purpose.)

6 Impact on the standard

The facilities proposed here are mostly pure extensions to the library except for minor changes to the `uses_allocator` trait and to types that use type erasure for allocators: `function`, `packaged_task`, `future`, `promise` and the upcoming `filepath` type in the file-system TS [N3399]. No core language changes are proposed.

7 Implementation Experience

The implementation of the new `memory_resource`, `resource_adaptor`, and `polymorphic_allocator` features is very straightforward. A prototype implementation based on this paper is available at http://www.halpernwightsoftware.com/WG21/polymorphic_allocator.tgz. The prototype also includes a rework of the gnu `function` class template to add the functionality described in this proposal. Most of the work in adapting `function` was in adding allocator support without breaking binary (ABI) compatibility.

The `memory_resource`, `polymorphic_allocator`, `monotonic_buffer_resource`, and `unsynchronized_pool_resource` classes described in this proposal are minor variations of the facilities that have been in use at Bloomberg for over a decade (See the BSL open-source library). These facilities have dramatically improved testability of software (through the use of test resources) and provided performance benefits when using special-purpose allocators such as arena allocators and thread-specific allocators.

8 Formal Wording – new classes

Throughout this wording, *Precondition* clauses are distinguished from *Requires* clauses. The former clause is used for run-time conditions that, if violated, would produce undefined behavior. The latter clause is used for compile-time conditions

8.1 Utility Class `erased_type`

Add a new section to the TS describing the following `erased_type` utility component:

u.1 Header `<experimental/utility>` synopsis [utility.syn]

```
#include <utility>

namespace std {
namespace experimental {

    // erased-type placeholder
    struct erased_type { };
```

Note to the editor: Other TS utility components from other papers would be merged into this synopsis here

```
    }
}
```

u.2 Class `erased_type` [erased.type]

```
namespace std {
namespace experimental {
    struct erased_type { };
}
}
```

The `erased_type struct` is an empty `struct` that serves as a placeholder for a type *T* in situations where the actual type *T* is determined at runtime. For example, the nested type, `allocator_type`, is an alias for `erased_type` in classes that use *type-erased allocators* (see [type.erased.allocator]).

Although the first (and currently only) use of `erased_type` is in the context of memory allocation, the concept of type erasure is not allocator-specific. Since there may be new uses for this type in the future, I elected to put it in `<utility>` instead of in `<memory>`.

8.2 Polymorphic Memory Resources

Add a new subsection in the TS for the polymorphic memory resources.

w.x Polymorphic Memory Resources [memory.resource]

8.2.1 Header `<experimental/memory_resource>` synopsis

w.x.1 Header `<experimental/memory_resource>` **synopsis [memory.resource.syn]**

```
namespace std {
```

```cpp
namespace experimental {
namespace pmr {

  class memory_resource;

  bool operator==(const memory_resource& a,
                  const memory_resource& b) noexcept;
  bool operator!=(const memory_resource& a,
                  const memory_resource& b) noexcept;

  template <class Tp> class polymorphic_allocator;

  template <class T1, class T2>
    bool operator==(const polymorphic_allocator<T1>& a,
                    const polymorphic_allocator<T2>& b) noexcept;
  template <class T1, class T2>
    bool operator!=(const polymorphic_allocator<T1>& a,
                    const polymorphic_allocator<T2>& b) noexcept;

  // The name resource_adaptor_imp is for exposition only.
  template <class Allocator> class resource_adaptor_imp;

  template <class Allocator>
    using resource_adaptor = resource_adaptor_imp<
      allocator_traits<Allocator>::rebind_alloc<char>>;

  // Global memory resources
  memory_resource* new_delete_resource() noexcept;
  memory_resource* null_memory_resource() noexcept;

  // The default memory resource
  memory_resource* set_default_resource(memory_resource* r) noexcept;
  memory_resource* get_default_resource() noexcept;

  // Standard memory resources
  struct pool_options;
  class synchronized_pool_resource;
  class unsynchronized_pool_resource;
  class monotonic_buffer_resource;

} // namespace pmr
} // namespace experimental
} // namespace std
```

8.2.2 Class `memory_resource`

w.x.2 Class `memory_resource` **[memory.resource.class]**

w.x.2.1 Class `memory_resource` **overview [memory.resource.overview]**

The `memory_resource` class is an abstract interface to an unbounded set of classes encapsulating memory resources.

```
namespace std {
namespace experimental {
namespace pmr {

  class memory_resource
  {
      // For exposition only
      static constexpr size_t max_align = alignof(max_align_t);

    public:
      virtual ~memory_resource();

      void* allocate(size_t bytes, size_t alignment = max_align);
      void deallocate(void* p, size_t bytes,
                      size_t alignment = max_align);

      bool is_equal(const memory_resource& other) const noexcept;

    protected:
      virtual void* do_allocate(size_t bytes, size_t alignment) = 0;
      virtual void  do_deallocate(void* p, size_t bytes,
                                  size_t alignment) = 0;

      virtual bool do_is_equal(const memory_resource& other) const
        noexcept = 0;
  };

} // namespace pmr
} // namespace experimental
} // namespace std
```

The use of the pattern whereby a public non-virtual function calls a protected virtual function enables default arguments to be expressed only once, in the abstract base class. It has the additional benefit of allowing an implementation to instrument the function or for the meaning of the function to evolve in the standard without breaking existing derived classes. Finally, this pattern is convenient for specification because it separates the public interface from the derived-class requirements.

w.x.2.2 `memory_resource` **public member functions [memory.resource.public]**

```
~memory_resource();
```

> *Effects:* Destroys this `memory_resource`.

```
void* allocate(size_t bytes, size_t alignment = max_align);
```

Effects: equivalent to `return do_allocate(bytes, alignment);`

`void deallocate(void* p, size_t bytes, size_t alignment = max_align);`

> *Effects:* equivalent to `do_deallocate(p, bytes, alignment);`

`bool is_equal(const memory_resource& other) const noexcept;`

> *Effects:* equivalent to `return do_is_equal(other);`

w.x.2.3 `memory_resource` protected virtual member functions [memory.resource.priv]

`virtual void* do_allocate(size_t bytes, size_t alignment) = 0;`

> *Preconditions:* `alignment` shall be a power of two.

> *Returns:* A derived class shall implement this function to return a pointer to allocated storage (3.7.4.2) with a size of at least `bytes`. The returned storage is aligned to the specified `alignment`, if such alignment is supported; otherwise it is aligned to `max_align`.

> [Note to editor: 3.7.4.2 [basic.stc.dynamic.deallocation]does not seem to actually define *allocated storage*, even though it is referenced in 3.8 [basic.life]. I could not find an actual definition of this term, but from the usage, it seems to mean storage that does not currently have an object constructed in it.]

> *Throws:* a derived class implementation shall throw an appropriate exception if it is unable to allocate memory with the requested size and alignment.

`virtual void do_deallocate(void* p, size_t bytes, size_t alignment) = 0;`

> *Preconditions:* `p` shall have been returned from a prior call to `allocate(bytes, alignment)` on a memory resource equal to `*this`, and the storage at `p` shall not yet have been deallocated.

> *Effects:* A derived class shall implement this function to dispose of allocated storage.

> *Throws:* Nothing.

Although this function throws nothing, it is not declared `noexcept` because it has a narrow interface. An implementation may choose to throw if a defensive test of the preconditions fails.

> `virtual bool do_is_equal(const memory_resource& other) const noexcept = 0;`

> > *Returns:* A derived class shall implement this function to return `true` if memory allocated from `this` can be deallocated from `other` and vice-versa; otherwise it shall return `false`. [*Note:* The most-derived type of `other` might not match the type of `this`. For a derived class, D, a typical implementation of this function will compute `dynamic_cast<D*>(&other)` and go no further (i.e., return `false`) if it returns `nullptr`. — *end note*]

For most classes derived from `memory_resource`, `do_is_equal` will return exactly `this == &other`. I.e., most memory resources are equal only if they are the same object. The `resource_adaptor` template (below) is a rare exception.

w.x.2.4 `memory_resource` equality [memory.resource.eq]

`bool operator==(const memory_resource& a,`
`              const memory_resource& b) noexcept;`

> *Returns:* `&a == &b || a.is_equal(b)`.

The explicit optimization of testing for `&a == &b` means that the implementation shall not invoke `is_equal` if the pointers compare equal. If this test were not explicit, then this important optimization would actually be illegal because the number of calls to `is_equal` is user-detectible.

```
bool operator!=(const memory_resource& a,
                const memory_resource& b) noexcept;
```

Returns: `!(a == b)`.

8.2.3 Class template `polymorphic_allocator`

w.x.3 Class template `polymorphic_allocator` [polymorphic.allocator.class]

w.x.3.1 Class template `polymorphic_allocator` overview [polymorphic.allocator.overview]

A specialization of class template `pmr::polymorphic_allocator` conforms to the Allocator requirements ([allocator.requirements] 17.6.3.5). Constructed with different memory resources, different instances of the same specialization of `pmr::polymorphic_allocator` can exhibit entirely different allocation behavior. This runtime polymorphism allows objects that use `polymorphic_allocator` to behave as if they used different allocator types at run time even though they use the same static allocator type.

```
namespace std {
namespace experimental {
namespace pmr {

  template <class Tp>
  class polymorphic_allocator
  {
      memory_resource* m_resource;    // For exposition only

    public:
      typedef Tp value_type;

      polymorphic_allocator() noexcept;
      polymorphic_allocator(memory_resource* r);

      polymorphic_allocator(const polymorphic_allocator& other)
        = default;

      template <class U>
        polymorphic_allocator(const polymorphic_allocator<U>& other)
          noexcept;

      polymorphic_allocator&
        operator=(const polymorphic_allocator& rhs) = default;

      Tp* allocate(size_t n);
      void deallocate(Tp* p, size_t n);

      template <typename T, typename... Args>
        void construct(T* p, Args&&... args);
```

```
                // Specializations for pair using piecewise construction
                template <class T1, class T2, class... Args1, class... Args2>
                  void construct(pair<T1,T2>* p, piecewise_construct_t,
                                 tuple<Args1...> x, tuple<Args2...> y);
                template <class T1, class T2>
                  void construct(pair<T1,T2>* p);
                template <class T1, class T2, class U, class V>
                  void construct(pair<T1,T2>* p, U&& x, V&& y);
                template <class T1, class T2, class U, class V>
                  void construct(pair<T1,T2>* p,
                                 const std::pair<U, V>& pr);
                template <class T1, class T2, class U, class V>
                  void construct(pair<T1,T2>* p, pair<U, V>&& pr);

                template <typename T>
                  void destroy(T* p);

                // Return a default-constructed allocator (no allocator propagation)
                polymorphic_allocator select_on_container_copy_construction()
                  const;

                memory_resource* resource() const;
            };

      } // namespace pmr
      } // namespace experimental
      } // namespace std
```

w.x.3.2 `polymorphic_allocator` constructors [polymorphic.allocator.ctor]

```
polymorphic_allocator() noexcept;
```

> *Effects:* Sets `m_resource` to `get_default_resource()`.

```
polymorphic_allocator(memory_resource* r);
```

> *Precondition:* `r` is non-null.

> *Effects:* Sets `m_resource` to `r`.

> *Throws:* nothing

> *Note:* This constructor provides an implicit conversion from `memory_resource*`.

```
template <class U>
  polymorphic_allocator(const polymorphic_allocator<U>& other) noexcept;
```

> *Effects:* sets `m_resource` to `other.resource()`.

w.x.3.3 `polymorphic_allocator` member functions [polymorphic.allocator.mem]

```
Tp* allocate(size_t n);
```

> *Returns:* Equivalent to `static_cast<Tp*>(m_resource->allocate(n * sizeof(Tp), alignof(Tp)))`.

```
void deallocate(Tp* p, size_t n);
```

Preconditions: p was allocated from a memory resource, *x*, equal to `*m_resource`, using `x.allocate(n * sizeof(Tp), alignof(Tp))`.

Effects: Equivalent to `m_resource->deallocate(p, n * sizeof(Tp), alignof(Tp))`.

Throws: Nothing.

```
template <class T, class... Args>
  void construct(T* p, Args&&... args);
```

Requires: uses-allocator construction of T with allocator `this->resource()` (see [mods.allocator.uses]) and constructor arguments `std::forward<Args>(args)...` is well-formed. [*Note:* uses-allocator construction is always well formed for types that do not use allocators. – *end note*]

Effects: Construct a `T` object at p by *uses-allocator construction* with allocator `this->resource()` ([mods.allocator.uses]) and constructor arguments `std::forward<Args>(args)...`.

Throws: Nothing unless the constructor for `T` throws.

```
template <class T1, class T2, class... Args1, class... Args2>
  void construct(pair<T1,T2>* p, piecewise_construct_t,
                 tuple<Args1...> x, tuple<Args2...> y);
```

Effects: Let `xprime` be a `tuple` constructed from x according to the appropriate rule from the following list. [*Note:* The following description can be summarized as constructing a `std::pair<T1,T2>` object at p as if by separate *uses-allocator construction* with allocator `this->resource()` ([allocator.uses.construction] 20.6.7.2) of `p->first` using the elements of x and `p->second` using the elements of y. – *end note*]:

— If `uses_allocator<T1,memory_resource*>::value` is `false` and `is_constructible<T,Args1...>::value` is `true`, then `xprime` is x.

— Otherwise, if `(uses_allocator<T1,memory_resource*>::value` is `true` and `is_constructible<T1,allocator_arg_t,memory_resource*,Args1...>::value)` is `true`, then `xprime` is `tuple_cat(make_tuple(allocator_arg, this->resource()), move(x))`.

— Otherwise, if `(uses_allocator<T1,memory_resource*>::value` is `true` and `is_constructible<T1,Args1...,memory_resource*>::value)` is `true`, then `xprime` is `tuple_cat(move(x), make_tuple(this->resource()))`.

— Otherwise the program is ill formed.

and let `yprime` be a `tuple` constructed from y according to the appropriate rule from the following list:

— If `uses_allocator<T2,memory_resource*>::value` is `false` and `is_constructible<T,Args2...>::value` is `true`, then `yprime` is y.

— Otherwise, if `(uses_allocator<T2,memory_resource*>::value` is `true` and `is_constructible<T2,allocator_arg_t,memory_resource*,Args2...>::value)` is `true`, then `yprime` is `tuple_cat(make_tuple(allocator_arg, this->resource()), move(y))`.

— Otherwise, if (`uses_allocator<T2,memory_resource*>::value` is true and `is_constructible<T2,Args2...,memory_resource*>::value`) is true, then `yprime` is `tuple_cat(move(y), make_tuple(this->resource()))`.

— Otherwise the program is ill formed.

then this function constructs a `std::pair<T1,T2>` object at `p` using constructor arguments `piecewise_construct, xprime, yprime`.

> The description above is almost identical to that in `scoped_allocator_adaptor` because a `polymorphic_allocator` is scoped. It differs in that, instead of passing `*this` down to the constructed object, it passes `this->resource()`.

```
template <class T1, class T2>
  void construct(std::pair<T1,T2>* p);
```

Effects: equivalent to `this->construct(p, piecewise_construct, tuple<>(), tuple<>())`;

```
template <class T1, class T2, class U, class V>
  void construct(std::pair<T1,T2>* p, U&& x, V&& y);
```

Effects: equivalent to `this->construct(p, piecewise_construct, forward_as_tuple(std::forward<U>(x)), forward_as_tuple(std::forward<V>(y)))`;

```
template <class T1, class T2, class U, class V>
  void construct(std::pair<T1,T2>* p, const std::pair<U, V>& pr);
```

Effects: equivalent to `this->construct(p, piecewise_construct, forward_as_tuple(pr.first), forward_as_tuple(pr.second))`;

```
template <class T1, class T2, class U, class V>
  void construct(std::pair<T1,T2>* p, std::pair<U, V>&& pr);
```

Effects: equivalent to `this->construct(p, piecewise_construct, forward_as_tuple(std::forward<U>(pr.first)), forward_as_tuple(std::forward<V>(pr.second)))`;

```
template <typename T>
  void destroy(T* p);
```

Effects: `p->~T()`.

```
polymorphic_allocator select_on_container_copy_construction() const;
```

Returns: `polymorphic_allocator()`.

```
memory_resource* resource() const;
```

Returns: `m_resource`.

w.x.3.4 `polymorphic_allocator` equality [polymorphic.allocator.eq]

```
template <class T1, class T2>
  bool operator==(const polymorphic_allocator<T1>& a,
                  const polymorphic_allocator<T2>& b) noexcept;
```

Returns: `*a.resource() == *b.resource()`.

```
template <class T1, class T2>
```

```
  bool operator!=(const polymorphic_allocator<T1>& a,
                  const polymorphic_allocator<T2>& b) noexcept;
```

  Returns: ! (a == b)

8.2.4   Class-alias template `resource_adaptor`

w.x.4 template alias `resource_adaptor` [resource.adaptor]

w.x.4.1 `resource_adaptor` [resource.adaptor.overview]

An instance of `resource_adaptor<Allocator>` is an adaptor that wraps a `memory_resource` interface around `Allocator`. In order that `resource_adaptor<X<T>>` and `resource_adaptor<X<U>>` are the same type for any allocator template X and types T and U, `resource_adaptor<Allocator>` is rendered as an alias to a class template such that `Allocator` is rebound to a `char` value type in every specialization of the class template. The requirements on this class template are defined below. The name *resource_adaptor_imp* is for exposition only and is not normative, but the definition of the members of that class, whatever its name, *are* normative.

In addition to the `Allocator` requirements ([allocator.requirements] 17.6.3.4), the parameter to `resource_adaptor` shall meet the following additional requirements:

- `typename allocator_traits<Allocator>::pointer` shall be identical to `typename allocator_traits<Allocator>::value_type*`.

- `typename allocator_traits<Allocator>::const_pointer` shall be identical to `typename allocator_traits<Allocator>::value_type const*`.

- `typename allocator_traits<Allocator>::void_pointer` shall be identical to `void*`.

- `typename allocator_traits<Allocator>::const_void_pointer` shall be identical to `void const*`.

```
namespace std {
namespace experimental {
namespace pmr {

  // The name resource_adaptor_imp is for exposition only.
  template <class Allocator>
    class resource_adaptor_imp : public memory_resource {

    // for exposition only
    Allocator m_alloc;

   public:
    typedef Allocator allocator_type;

    resource_adaptor_imp() = default;
    resource_adaptor_imp(const resource_adaptor_imp&) = default;
    resource_adaptor_imp(resource_adaptor_imp&&) = default;

    explicit resource_adaptor_imp(const Allocator& a2);
    explicit resource_adaptor_imp(Allocator&& a2);
```

```
          resource_adaptor_imp& operator=(const resource_adaptor_imp&)
            = default;

          allocator_type get_allocator() const { return m_alloc; }

        protected:
          virtual void* do_allocate(size_t bytes, size_t alignment);
          virtual void  do_deallocate(void* p, size_t bytes,
                                      size_t alignment);

          virtual bool do_is_equal(const memory_resource& other) const
            noexcept;
        };

    template <class Allocator>
      using resource_adaptor = typename resource_adaptor_imp<
        allocator_traits<Allocator>::template rebind_alloc<char>>;

  } // namespace pmr
  } // namespace experimental
  } // namespace std
```

w.x.4.2 `resource_adaptor_imp` constructors [resource.adaptor.ctor]

`explicit` *`resource_adaptor_imp`*`(const Allocator& a2);`

> *Effects:* Initializes `m_alloc` with `a2`.

`explicit` *`resource_adaptor_imp`*`(Allocator&& a2);`

> *Effects:* Initializes `m_alloc` with `std::move(a2)`.

w.x.4.3 `resource_adaptor_imp` member functions [resource.adaptor.mem]

`void* do_allocate(size_t bytes, size_t alignment);`

> *Returns:* Allocated memory obtained by calling `m_alloc.allocate`. The size and alignment of the allocated memory shall meet the requirements for a class derived from `memory_resource` ([memory.resource]).

`void do_deallocate(void* p, size_t bytes, size_t alignment);`

> *Requires:* `p` was previously allocated using A`.allocate`, where A `== m_alloc`, and not subsequently deallocated.

> *Effects:* Returns memory to the allocator using `m_alloc.deallocate()`.

`bool do_is_equal(const memory_resource& other) const;`

> Let `p` be `dynamic_cast<const` *`resource_adaptor_imp`*`*>(&other)`.

> *Returns:* `false` if `p` is null, otherwise the value of `m_alloc == p->m_alloc`.

8.2.5   Program-wide `memory_resource` objects

w.x.5 Access to program-wide `memory_resource` objects [memory.resource.global]

```
memory_resource* new_delete_resource() noexcept;
```

> *Returns:* A pointer to a static-duration object of a type derived from `memory_resource` that can serve as a resource for allocating memory using `::operator new` and `::operator delete`. The same value is returned every time this function is called. For return value `p` and memory resource `r`, `p->is_equal(r)` returns `&r == p`.

```
memory_resource* null_memory_resource() noexcept;
```

> *Returns:* A pointer to a static-duration object of a type derived from `memory_resource` for which `allocate()` always throws `bad_alloc` and for which `deallocate()` has no effect. The same value is returned every time this function is called. For return value `p` and memory resource `r`, `p->is_equal(r)` returns `&r == p`.

A memory resource may obtain memory using another resource for replenishing its pool. The null memory resource is useful for situations where the original pool is not expected to become exhausted.

The *default memory resource pointer* is a pointer to a memory resource that is used by certain facilities when an explicit memory resource is not supplied through the interface. Its initial value is the return value of `new_delete_resource()`.

```
memory_resource* set_default_resource(memory_resource* r) noexcept;
```

> *Effects:* If `r` is non-null, sets the value of the default memory resource pointer to `r`, otherwise sets the default memory resource pointer to `new_delete_resource()`.

> *Post-condition:* `get_default_resource() == r`.

> *Returns:* The previous value of the default memory resource pointer.

> *Remarks:* Calling the `set_default_resource` and `get_default_resource` functions shall not incur a data race. A call to the `set_default_resource` function shall synchronize with subsequent calls to the `set_default_resource` and `get_default_resource` functions.

These synchronization requirements are the same as for `set/get_new_handler` and `set/get_terminate`.

```
memory_resource* get_default_resource() noexcept;
```

> *Returns:* The current value of the default memory resource pointer.

8.3   Classes `synchronized_pool_resource` and `unsynchronized_pool_resource`

w.x.6 Pool resource classes [pool.resource]

w.x.6.1 Classes `synchronized_pool_resource` and `unsynchronized_pool_resource` [pool.resource.overview]

The `synchronized_pool_resource` and `unsynchronized_pool_resource` classes (collectively, *pool resource* classes) are general-purpose memory resources having the following qualities:

— Each resource *owns* the allocated memory, and frees it on destruction – even if `deallocate` has not been called for some of the allocated blocks.

— A pool resource (see Figure 1) consists of a collection of pools, serving requests for different block sizes. Each individual pool manages a collection of *chunks* that are in turn divided into blocks of uniform size, returned via calls to `do_allocate`. Each call to `do_allocate(size, alignment)` is dispatched to the pool serving the smallest blocks accommodating at least `size` bytes.

— When a particular pool is exhausted, allocating a block from that pool results in the allocation of an additional chunk of memory from the *upstream allocator* (supplied at construction), thus replenishing the pool. With each successive replenishment, the chunk size obtained increases geometrically. [*Note:* By allocating memory in chunks, the pooling strategy increases the chance that consecutive allocations will be close together in memory. *– end note*]

— Allocation requests that exceed the largest block size of any pool are fulfilled directly from the upstream allocator.

— A `pool_options` struct may be passed to the pool resource constructors to tune the largest block size and the maximum chunk size.

[*Example:* Figure 1 shows a possible data structure that implements a pool resource.

Figure 1: pool resource

– end example]

A `synchronized_pool_resource` may be accessed from multiple threads without external synchronization and may have thread-specific pools to reduce synchronization costs. An `unsynchronized_pool_resource` class may not be accessed from multiple threads simultaneously and thus avoids the cost of synchronization entirely in single-threaded applications.

```
namespace std {
namespace experimental {
namespace pmr {
```

```
struct pool_options
{
  size_t              max_blocks_per_chunk = 0;
  size_t              largest_required_pool_block = 0;
};
```

By bundling the options together into a `struct`, the interface is substantially more future-proof than if the options were specified individually in the constructors. It is much easier to add (named) fields to a `struct` than to overload on a list of (unnamed) arguments to a constructor.

```
class synchronized_pool_resource : public memory_resource
{
  public:
    synchronized_pool_resource(const pool_options& opts,
                               memory_resource* upstream);

    synchronized_pool_resource()
      : synchronized_pool_resource(pool_options(),
                                   get_default_resource()) { }
    explicit synchronized_pool_resource(memory_resource* upstream)
      : synchronized_pool_resource(pool_options(), upstream) { }
    explicit synchronized_pool_resource(pool_options& opts)
      : synchronized_pool_resource(opts,
                                   get_default_resource()) { }

    synchronized_pool_resource(
      const synchronized_pool_resource&) = delete;
    virtual ~synchronized_pool_resource();

    synchronized_pool_resource& operator=(
      const synchronized_pool_resource&) = delete;

    void release();
    memory_resource* upstream_resource() const;
    pool_options     options() const;

  protected:
    virtual void* do_allocate(size_t bytes, size_t alignment);
    virtual void  do_deallocate(void* p, size_t bytes,
                                size_t alignment);

    virtual bool do_is_equal(const memory_resource& other) const
      noexcept;
};

class unsynchronized_pool_resource : public memory_resource
{
  public:
    unsynchronized_pool_resource(const pool_options& opts,
```

```
                                      memory_resource* upstream);

            unsynchronized_pool_resource()
              : unsynchronized_pool_resource(pool_options(),
                                             get_default_resource()) { }
            explicit unsynchronized_pool_resource(memory_resource* upstream)
              : unsynchronized_pool_resource(pool_options(), upstream) { }
            explicit unsynchronized_pool_resource(const pool_options& opts)
              : unsynchronized_pool_resource(opts,
                                             get_default_resource()) { }

            unsynchronized_pool_resource(
              const unsynchronized_pool_resource&) = delete;
            virtual ~unsynchronized_pool_resource();

            unsynchronized_pool_resource& operator=(
              const unsynchronized_pool_resource&) = delete;

            void release();
            memory_resource* upstream_resource() const;
            pool_options     options() const;

          protected:
            virtual void* do_allocate(size_t bytes, size_t alignment);
            virtual void  do_deallocate(void* p, size_t bytes,
                                        size_t alignment);

            virtual bool do_is_equal(const memory_resource& other) const
              noexcept;
        };

    } // namespace pmr
    } // namespace experimental
    } // namespace std
```

w.x.6.2 `pool_options` data members

The members of `pool_options` comprise a set of constructor options for pool resources. The effect of each option on the pool resource behavior is described below:

`size_t max_blocks_per_chunk;`

> The maximum number of blocks that will be allocated at once from the upstream memory resource to replenish a pool.

> If the value of `max_blocks_per_chunk` is zero or is greater than an implementation-defined limit, that limit is used instead. The implementation may choose to use a smaller value than is specified in this field and may use different values for different pools.

`size_t largest_required_pool_block;`

> The largest allocation size that is required to be fulfilled using the pooling mechanism. Attempts to allocate a single block larger than this threshold will be allocated directly from the upstream memory

resource. If `largest_required_pool_block` is zero or is greater than an implementation-defined limit, that limit is used instead. The implementation may choose a pass-through threshold larger than specified in this field.

w.x.6.3 pool resource constructors and destructors [pool.ctor]

```
synchronized_pool_resource(const pool_options& opts,
                           memory_resource* upstream);
unsynchronized_pool_resource(const pool_options& opts,
                             memory_resource* upstream);
```

Precondition: `upstream` is the address of a valid memory resource.

Effects: Constructs a pool resource object that will obtain memory from `upstream` whenever the pool resource is unable to satisfy a memory request from its own internal data structures. The resulting object will hold a copy of `upstream`, but will not own the resource to which `upstream` points. [*Note:* The intention is that calls to `upstream->allocate()` will be substantially fewer than calls to `this->allocate()` in most cases. – *end note*] The behavior of the pooling mechanism is tuned according to the value of the `opts` argument.

Throws: Nothing unless `upstream->allocate()` throws. It is unspecified if or under what conditions this constructor calls `upstream->allocate()`.

```
virtual ~synchronized_pool_resource();
virtual ~unsynchronized_pool_resource();
```

Effects: calls `this->release()`.

w.x.6.4 pool resource members [pool.mem]

```
void release();
```

Effects: Calls `upstream_resource()->deallocate()` as necessary to release all allocated memory. [*Note:* memory is released back to `upstream_resource()` even if `deallocate` has not been called for some of the allocated blocks. – *end note*]

```
memory_resource* upstream_resource() const;
```

Returns: The value of the `upstream` argument provided to the constructor of this object.

```
pool_options options() const;
```

Returns: The options that control the pooling behavior of this resource. The values in the returned `struct` may differ from those supplied to the pool resource constructor in that values of zero will be replaced with implementation-defined defaults and sizes may be rounded to unspecified granularity.

```
virtual void* do_allocate(size_t bytes, size_t alignment);
```

Returns: A pointer to allocated storage (3.7.4.2) with a size of at least `bytes`. The size and alignment of the allocated memory shall meet the requirements for a class derived from `memory_resource` ([memory.resource]).

Effects: If the pool selected for a block of size `bytes` is unable to satisfy the memory request from its own internal data structures, it will call `upstream_resource()->allocate()` to obtain more memory. If `bytes` is larger than that which the largest pool can handle, then memory will be allocated using `upstream_resource()->allocate()`.

Throws: Nothing unless `upstream_resource()->allocate()` throws.

```
virtual void do_deallocate(void* p, size_t bytes, size_t alignment);
```

> *Effects:* Return the memory at `p` to the pool. It is unspecified if or under what circumstances this operation will result in a call to `upstream_resource()->deallocate()`.

> *Throws:* Nothing

```
virtual bool
unsynchronized_pool_resource::do_is_equal(const memory_resource& other)
    const noexcept;
```

> *Returns:* `this == dynamic_cast<unsynchronized_pool_resource*>(&other)`.

```
virtual bool
synchronized_pool_resource::do_is_equal(const memory_resource& other)
    const noexcept;
```

> *Returns:* `this == dynamic_cast<synchronized_pool_resource*>(&other)`.

8.4 Class `monotonic_buffer_resource`

w.x.7 Class `monotonic_buffer_resource` [monotonic.buffer]

w.x.7.1 Class `monotonic_buffer_resource` overview [monotonic.buffer]

A `monotonic_buffer_resource` is a special-purpose memory resource intended for very fast memory allocations in situations where memory is used to build up a few objects and then is released all at once when the memory resource object is destroyed. It has the following qualities:

— A call to `deallocate` has no effect, thus the amount of memory consumed increases monotonically until the resource is destroyed.

— The program can supply an initial buffer, which the allocator uses to satisfy memory requests.

— When the initial buffer (if any) is exhausted, it obtains additional buffers from an *upstream* memory resource supplied at construction. Each additional buffer is larger than the previous one, following a geometric progression.

— It is intended for access from one thread of control at a time. Specifically, calls to `allocate` and `deallocate` do not synchronize with one another.

— It *owns* the allocated memory and frees it on destruction, even if `deallocate` has not been called for some of the allocated blocks.

```
namespace std {
namespace experimental {
namespace pmr {

  class monotonic_buffer_resource : public memory_resource
  {
      memory_resource* upstream_rsrc;      // exposition only
      void*            current_buffer;     // exposition only
      size_t           next_buffer_size;   // exposition only

    public:
      explicit monotonic_buffer_resource(memory_resource* upstream);
      monotonic_buffer_resource(size_t initial_size,
                                memory_resource* upstream);
```

```cpp
      monotonic_buffer_resource(void* buffer, size_t buffer_size,
                                memory_resource* upstream);

      monotonic_buffer_resource()
        : monotonic_buffer_resource(get_default_resource()) { }
      explicit monotonic_buffer_resource(size_t initial_size)
        : monotonic_buffer_resource(initial_size,
                                    get_default_resource()) { }
      monotonic_buffer_resource(void* buffer, size_t buffer_size)
        : monotonic_buffer_resource(buffer, buffer_size,
                                    get_default_resource()) { }

      monotonic_buffer_resource(const monotonic_buffer_resource&)
        = delete;

      virtual ~monotonic_buffer_resource();

      monotonic_buffer_resource operator=(
        const monotonic_buffer_resource&) = delete;

      void release();
      memory_resource* upstream_resource() const;

    protected:
      virtual void* do_allocate(size_t bytes, size_t alignment);
      virtual void  do_deallocate(void* p, size_t bytes,
                                  size_t alignment);

      virtual bool do_is_equal(const memory_resource& other) const
        noexcept;
    };

  } // namespace pmr
  } // namespace experimental
  } // namespace std
```

w.x.7.2 `monotonic_buffer_resource` constructor and destructor [monotonic.buffer.ctor]

```cpp
explicit monotonic_buffer_resource(memory_resource* upstream);
monotonic_buffer_resource(size_t initial_size, memory_resource* upstream);
```

Preconditions: `upstream` shall be the address of a valid memory resource; `initial_size`, if specified, shall be greater than zero.

Effects: Sets `upstream_rsrc` to `upstream` and `current_buffer` to `nullptr`. If `initial_size` is specified, sets `next_buffer_size` to at least `initial_size`; otherwise sets `next_buffer_size` to an implementation-defined size.

```cpp
monotonic_buffer_resource(void* buffer, size_t buffer_size,
                          memory_resource* upstream);
```

Preconditions: `upstream` shall be the address of a valid memory resource. `buffer_size` shall be no larger than the number of bytes in `buffer`.

Effects: Sets `upstream_rsrc` to `upstream`, `current_buffer` to `buffer`, and `next_buffer_size` to `initial_size` (but not less than 1), then increases `next_buffer_size` by an implementation-defined growth factor (which need not be integral).

```
~monotonic_buffer_resource();
```

Effects: Calls `this->release()`.

w.x.7.3 `monotonic_buffer_resource` members [monotonic.buffer.mem]

```
void release();
```

Effects: Calls `upstream_rsrc->deallocate()` as necessary to release all allocated memory. [*Note:* memory is released back to `upstream_rsrc` even if some blocks that were allocated from `this` have not been deallocated from `this`. – *end note*]

```
memory_resource* upstream_resource() const;
```

Returns: the value of `upstream_rsrc`.

```
void* do_allocate(size_t bytes, size_t alignment);
```

Returns: A pointer to allocated storage (3.7.4.2) with a size of at least `bytes`. The size and alignment of the allocated memory shall meet the requirements for a class derived from `memory_resource` ([memory.resource]).

Effects: If the unused space in `current_buffer` can fit a block with the specified `bytes` and `alignment`, then allocate the return block from `current_buffer`; otherwise set `current_buffer` to `upstream_rsrc->allocate(n, m)`, where n is not less than `max(bytes, next_buffer_size)` and m is not less than `alignment`, and increase `next_buffer_size` by an implementation-defined growth factor (which need not be integral), then allocate the return block from the newly-allocated `current_buffer`.

Throws: Nothing unless `upstream_rsrc->allocate()` throws.

```
void do_deallocate(void* p, size_t bytes, size_t alignment);
```

Effects: None

Throws: Nothing

Remarks: Memory used by this resource increases monotonically until its destruction.

```
bool do_is_equal(const memory_resource& other) const noexcept;
```

Returns: `this == dynamic_cast<monotonic_buffer_resource*>(&other)`.

8.5 String Aliases Using Polymorphic Allocators

Create an experimental extension to `<string>` to add variations of the standard string types that allocate memory using `pmr::polymorphic_allocator`:

w.x.8 Header `<experimental/string>` synopsis:

```
#include <string>

namespace std {
namespace experimental {
namespace pmr {
```

```
// basic_string using polymorphic allocator in namespace pmr
template <class charT, class traits = char_traits<charT>>
  using basic_string =
    std::basic_string<charT, traits, polymorphic_allocator<charT>>;


// basic_string typedef names using polymorphic allocator in namespace
// std::experimental::pmr
typedef basic_string<char> string;
typedef basic_string<char16_t> u16string;
typedef basic_string<char32_t> u32string;
typedef basic_string<wchar_t> wstring;


} // namespace std
} // namespace pmr
} // namespace experimental
```

With this change `pmr::wstring` is a `wstring` that uses a polymorphic allocator.

8.6 Containers Aliases Using Polymorphic Allocators

Create experimental extensions to most of the container headers to add variations of the standard containers that allocate memory using `pmr::polymorphic_allocator`:

w.x.9 Header `<experimental/deque>` synopsis [deque.syn]

```
#include <deque>

namespace std {
namespace experimental {
namespace pmr {

  template <class T>
    using deque = std::deque<T,polymorphic_allocator<T>>;

}
}
}
```

w.x.10 Header `<experimental/forward_list>` synopsis [forward_list.syn]

```
#include <forward_list>

namespace std {
namespace experimental {
namespace pmr {

  template <class T>
    using forward_list =
      std::forward_list<T,polymorphic_allocator<T>>;

}
}
}
```

w.x.11 Header **<experimental/list>** synopsis [list.syn]

```
#include <list>

namespace std {
namespace experimental {
namespace pmr {

  template <class T>
    using list = std::list<T,polymorphic_allocator<T>>;

}
}
}
```

w.x.12 Header **<experimental/vector>** synopsis [vector.syn]

```
#include <vector>

namespace std {
namespace experimental {
namespace pmr {

  template <class T>
    using vector = std::vector<T,polymorphic_allocator<T>>;

}
}
}
```

w.x.13 Header **<experimental/map>** synopsis [map.syn]

```
#include <map>

namespace std {
namespace experimental {
namespace pmr {

  template <class Key, class T, class Compare = less<Key>>
    using map = std::map<Key, T, Compare,
                polymorphic_allocator<pair<const Key,T>>>;

  template <class Key, class T, class Compare = less<Key>>
    using multimap = std::multimap<Key, T, Compare,
                polymorphic_allocator<pair<const Key,T>>>;

}
}
}
```

w.x.14 Header **<experimental/set>** synopsis [set.syn]

```
#include <set>

namespace std {
```

```
namespace experimental {
namespace pmr {

  template <class Key, class Compare = less<Key>>
    using set = std::set<Key, Compare,
                         polymorphic_allocator<Key>>;

  template <class Key, class Compare = less<Key>>
    using multiset = std::multiset<Key, Compare,
                         polymorphic_allocator<Key>>;

}
}
}
```

w.x.15 Header <experimental/unordered_map> synopsis [unordered_map.syn]

```
#include <unordered_map>

namespace std {
namespace experimental {
namespace pmr {

  template <class Key, class T,
            class Hash = hash<Key>,
            class Pred = std::equal_to<Key>>
    using unordered_map =
      std::unordered_map<Key, T, Hash, Pred,
                 polymorphic_allocator<pair<const Key,T>>>;

  template <class Key, class T,
            class Hash = hash<Key>,
            class Pred = std::equal_to<Key>>
    using unordered_multimap =
      std::unordered_multimap<Key, T, Hash, Pred,
                 polymorphic_allocator<pair<const Key,T>>>;

}
}
}
```

w.x.16 Header <experimental/unordered_set> synopsis [unordered_set.syn]

```
#include <unordered_set>

namespace std {
namespace experimental {
namespace pmr {

  template <class Key,
            class Hash = hash<Key>,
            class Pred = equal_to<Key>>
    using unordered_set =
```

```
            std::unordered_set<Key, Hash, Pred,
                                polymorphic_allocator<Key>>;

   template <class Key,
             class Hash = hash<Key>,
             class Pred = equal_to<Key>>
     using unordered_multiset =
       std::unordered_multiset<Key, Hash, Pred,
                               polymorphic_allocator<Key>>;
}
}
}
```

w.x.16 Header `<experimental/regex>` synopsis [regex.syn]

```
#include <regex>
#include <exerimental/string>

namespace std {
namespace experimental {
namespace pmr {

   template <class BidirectionalIterator>
     using match_results =
       std::match_results<BidirectionalIterator,
             polymorphic_allocator<sub_match<BidirectionalIterator>>>;

   typedef match_results<const char*> cmatch;
   typedef match_results<const wchar_t*> wcmatch;
   typedef match_results<string::const_iterator> smatch;
   typedef match_results<wstring::const_iterator> wsmatch;

}
}
}
```

8.7 Type-erased allocators

Insert a new section into the TS as follows:

x.y.z Type-erased allocator [type.erased.allocator]

A *type-erased allocator* is an allocator or memory resource, `alloc`, used to allocate internal data structures for an object X of type C, but where C is not dependent on the type of `alloc`. Once `alloc` has been supplied to X (typically as a constructor argument), `alloc` can be retrieved from X only as a pointer `rptr` of static type `std::experimental::pmr::memory_resource*` ([memory.resource.class]). The process by which `rptr` is computed from `alloc` depends on the type of `alloc` as described in Table Q:

Table Q – Computed `memory_resource` for type-erased allocator

If the type of `alloc` is	then the value of `rptr` is

If the type of `alloc` is	then the value of `rptr` is
non-existent – no `alloc` specified	The value of `experimental::pmr::get_default_resource()` at the time of construction.
`nullptr_t`	The value of `experimental::pmr::get_default_resource()` at the time of construction.
a pointer type convertible to `pmr::memory_resource*`	`static_cast<` `experimental::pmr::memory_resource*>(alloc)`
`pmr::polymorphic_allocator<U>`	`alloc.resource()`
any other type meeting the Allocator requirements ([allocator.requirements])	a pointer to a value of type `experimental::pmr::resource_adaptor<A>` where A is the type of `alloc`. `rptr` remains valid only for the lifetime of X.
None of the above	The program is ill-formed.

Additionally, class `C` shall meet the following requirements:

- `C::allocator_type` shall be identical to `std::experimental::erased_type`.

- `X.get_memory_resource()` returns `rptr`.

9 Formal wording – Changes to classes in the standard

Although this proposal is targeted towards a TS, there are a small number of C++14 standard library classes that would need to be adjusted in order for users of the TS to get maximum value from the features proposed here. These changes are expressed as deltas from the standard, but new classes are still within the experimental namespace. Implementers might wish have a macro to turn these new features on or off, depending on whether TS functionality is desired. (Such a macro could be added to the feature-test recommendations listed in N3694.)

Note: the section numbers below are relative to the October 2013 Committee Draft, N3797 and will need to be updated when the FDIS and eventually the IS is issued.

Add the following section to the TS:

t.u Departure from the ISO standard [standard.departure]

t.u.1 In general

Although most of the facilities described in this technical specification are strictly supplements to the ISO C++ Standard, a few facilities depend on extensions to entities within the standard itself. The following sections describe these extensions by quoting the affected parts of the standard and using underlining to represent added text and ~~strike-through~~ to represent deleted text.

The description above is a placeholder for text that might appear only once within the TS. I expect the project editors to edit or replace it as appropriate.

t.u.2 Uses-allocator construction [mods.allocator.uses]

The following changes to the `uses_allocator` trait and to the description of *uses-allocator construction* allow a `memory_resource` pointer act as an allocator in many circumstances. [*Note:* Existing programs that uses standard allocators would be unaffected by this change. – *end note*]

20.7.7 uses_allocator [allocator.uses]

20.7.7.1 uses_allocator trait [allocator.uses.trait]

```
template <class T, class Alloc> struct uses_allocator;
```

> *Remark*: automatically detects whether `T` has a nested `allocator_type` that is convertible from `Alloc`. Meets the BinaryTypeTrait requirements (20.9.1). The implementation shall provide a definition that is derived from `true_type` if a type `T::allocator_type` exists and either `is_convertible<Alloc, T::allocator_type>::value != false` or `T::allocator_type` is an alias for `std::experimental::erased_type` ([utility.erased_type]), otherwise it shall be derived from `false_type`. A program may specialize this template to derive from `true_type` for a user-defined type `T` that does not have a nested `allocator_type` but nonetheless can be constructed with an allocator where either:
>
> — the first argument of a constructor has type `allocator_arg_t` and the second argument has type `Alloc` or
>
> — the last argument of a constructor has type `Alloc`.

20.7.7.2 uses-allocator construction [allocator.uses.construction]

> *Uses-allocator construction with allocator* `Alloc` refers to the construction of an object `obj` of type `T`, using constructor arguments `v1, v2, ..., vN` of types `V1, V2, ..., VN`, respectively, and an allocator `alloc` of type `Alloc`, where `Alloc` either (1) meets the requirements of an allocator ([allocator.requirements]), or (2) is a pointer type convertible to `std::experimental::pmr::memory_resource*` ([polymorphic.allocator]), according to the following rules:

The new text for *Uses-allocator construction* is not strictly necessary, but it is intended to clarify that two different kinds of thing can be passed as `alloc` in uses-allocator construction.

9.1 Type-erased allocator for `function`

t.u.3 Additions to `std::function` [mods.func.wrap]

In section 20.9.11.2 [func.wrap.func], the following declarations are added as public members of class template `function`:

```
    typedef experimental::erased_type allocator_type;

    experimental::pmr::memory_resource* get_memory_resource();
```

In section 20.9.11.2.1 [func.wrap.func.con], the introductory paragraph is changed as follows, giving the constructors of the `function` class template support for a type-erased allocator:

> When a `function` constructor that takes a first argument of type `allocator_arg_t` is invoked, the second argument is treated as a *type-erased allocator* ([type.erased.allocator]). ~~shall have a type that conforms to the requirements for Allocator (Table 17.6.3.5). A copy of the allocator argument is used to allocate memory, if necessary, for the internal data structures of the constructed function object.~~ If the constructor moves or makes a copy of a function object (including an instance of the `function` class template), then that move or copy is performed by *using-allocator construction* with allocator `get_memory_resource()`.

In section 20.9.11.2.1 [func.wrap.func.con], the assignment operators are enhanced to take the type-erased allocator into account:

```
function& operator=(const function& f);
```

> *Effects*: function(allocator_arg, get_memory_resource(), f).swap(*this);

> *Returns*: *this

```
function& operator=(function&& f);
```

> *Effects*: ~~Replaces the target of *this with the target of f~~. function(allocator_arg, get_memory_resource(), std::move(f)).swap(*this);

> *Returns*: *this

```
function& operator=(nullptr_t);
```

> *Effects*: If *this != NULL, destroys the target of this.

> *Postconditions*: !(*this).

> *Returns*: *this

```
template<class F> function& operator=(F&& f);
```

> *Effects*: function(allocator_arg, get_memory_resource(), std::forward<F>(f)).swap(*this);

> *Returns*: *this

```
template<class F> function& operator=(reference_wrapper<F> f) ~~noexcept~~;
```

> *Effects*: function(allocator_arg, get_memory_resource(), f).swap(*this);

> *Returns*: *this

In section 20.9.11.2.2 [func.wrap.func.mod] a precondition is added to the definition of swap:

```
void swap(function& other) ~~noexcept~~;
```

> *Precondition*: this->get_memory_resource() == other->get_memory_resource().

> *Effects:* Interchanges the targets of *this and other.

9.2 Type-erased allocator for `promise`

t.u.4 Additions to `std::promise` [mods.futures.promise]

In section 30.6.5 [futures.promise], the following declarations are added as public members of class template promise:

```
typedef experimental::erased_type allocator_type;

experimental::pmr::memory_resource* get_memory_resource();
```

and the following paragraph is inserted before the first (introductory) paragraph of the section.

> When a promise constructor that takes a first argument of type allocator_arg_t is invoked, the second argument is treated as a *type-erased allocator* ([type.erased.allocator]).

9.3 Type-erased allocator for `packaged_task`

t.u.5 Additions to `std::packaged_task` [mods.futures.task]

In section 30.6.9 [futures.task], the following declarations are added as public members of class template `packaged_task`:

```
typedef experimental::erased_type allocator_type;

experimental::pmr::memory_resource* get_memory_resource();
```

and the following paragraph is inserted before the first (introductory) paragraph of the section.

When a `packaged_task` constructor that takes a first argument of type `allocator_arg_t` is invoked, the second argument is treated as a *type-erased allocator* ([type.erased.allocator]).

10 Appendix: Template Implementation Policy (Section 4.3 from N1850)

The first problem most people see with the allocator mechanism as specified in the Standard is that the choice of allocator affects the type of a container. Consider, for example, the following type and object definitions:

```
typedef std::list<int, std::allocator<int> > NormIntList;
typedef std::list<int, MyAllocator<int> >    MyIntList;

NormIntList list1(5, 3);
MyIntList   list2(5, 3);
```

`list1` and `list2` are both lists of integers, and both contain five copies of the number 3. Most people would say that they have the same *value*. Yet they belong to different types and you cannot substitute one for the other. For example, assume we have a function that builds up a list:

```
int build(std::list<int>& theList);
```

Because we did not specify an allocator parameter for the argument type, the default, `std::allocator<int>` is used. Thus, `theList` is a reference to the same type as `list1`. We can use `build` to put values into `list1`, but we cannot use it to put values into `list2` because `MyIntList` is not compatible with `std::list<int>`. The following operations are also not supported:

```
list1 == list2
list1 = list2
MyIntList list3(list1);
NormIntList* p = &list2;
// etc.
```

Now, some would argue that the solution to the `build` function problem is to templatize `build`:

```
template <typename Alloc>
int build(std::list<int, Alloc>& theList);
```

or, better yet:

```
template <typename OutputIterator>
int build(OutputIterator theIter);
```

Both of these templatized solutions have their place, but both add substantial complexity to the development process. Templates, if overused, lead to long compile times and, sometimes, bloated code. If `build` were a template and passed its arguments on to other functions, those functions would also need to be templates. This chained instantiation of templates produces a deep compile-time dependency such that a change to any of those modules would result in a recompilation of a significant part of the system. For thorough coverage of the benefits of reducing physical dependencies, see [Lakos96].

Even if the templatization solution were acceptable, once a nested container (e.g. a list of strings) is involved, even the simplest operations require many layers of code to bridge the type-interoperablity gap. Consider trying to compare a shared list of shared strings with a regular list of regular strings:

```
typedef std::basic_string<
        char,
        std::char_traits<char>,
        shared_alloc<char>
    > shared_string;

std::list<shared_string, shared_alloc<shared_string> > SharedList;
std::list<std::string> TestList;
```

Not only will `SharedList == TestList` fail to compile, but employing iterators and standard algorithms will not work either:

```
bool same = std::range_equal(SharedList.begin(), SharedList.end(),
                             TestList.begin(), TestList.end());
```

The types to which the iterators refer are not equality-compatible (`std::string` vs. `shared_string`). The interoperability barrier caused by the use of template implementation policies impedes the straightforward use of *vocabulary types* – ubiquitous types used throughout the internal interfaces of a program. For example, to declare a string, `s` using `MyAllocator` we would need to write

```
std::basic_string<char, std::char_traits<char>, MyAllocator<char> > s;
```

Many people find this hard to read, but the more important fact is that `s` is not an `std::string` object and cannot be used wherever `std::string` is expected. Similar problems exist for other common types like `std::vector<int>`. The use of a well-defined set of vocabulary types like `string` and `vector` lends simplicity and clarity to a piece of code. Unfortunately, their use hinders the effective use of STL-style allocators and vice-versa.

Finally, template code is much harder to test than non-template code. Templates do not produce executable machine code until instantiated. Since there are an unbounded number of possible instantiations for any given template, the number of test cases needed to ensure that every path is covered can grow by an order of magnitude for each template parameter. Subtle assumptions that the template writer makes about the template's parameters may not become apparent until someone instantiates the template with an innocent-looking, but not-quite-

compatible parameter, long after the engineer who created the template has left the project.

Template implementation policies can be very useful when constructing mechanisms, as in the case of a function object (functor) type being used to specify an implementation policy for a standard algorithm template. Alexandrescu makes a compelling case for the use of template class policies in situations where instantiations are not expected to interoperate. However, template implementation policies are detrimental when used to control the memory allocation mechanisms of basic types that could otherwise interoperate.

11 Acknowledgements

I'd like to thank John Lakos for his careful review of my introductory text and for showing me what allocators can really do, if correctly conceived. Also, a big thank you to the members of my former team at Bloomberg for your help in defining the concepts in this paper and reviewing the result, especially Alexander Beels, Henry Mike Verschell, and Alisdair Merideth, who reworked the usage example for me. Thanks to Mark Boyall for promoting the addition of new allocators to the standard and for reviewing an early draft of this paper.

12 References

N3726 *Polymorphic Memory Resources,* Pablo Halpern, 2013

N3525 *Polymorphic Allocators*, Pablo Halpern, 2013

N3575 *Additional Standard allocation schemes*, Mark Boyall, 2013

N1850 *Towards a Better Allocator Model*, Pablo Halpern, 2005

BSL   *The BSL open-source library,* BLoomberg

jsmith *C++ Type Erasure*, JSmith, Published on www.cplusplus.org, 2010-01-27

N3399 *Filesystem Library Proposal (Revision 3),* Beman Dawes, 2012-09-21